

Exploratory Data Analysis for 2-18 COD data

Load settings and data

```
source("code/00_config.R")

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.0      v purrr  0.3.4
## v tibble  3.1.8      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
##
## Attaching package: 'lubridate'
##
##
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
##
## Linking to GEOS 3.10.2, GDAL 3.4.2, PROJ 8.2.1; sf_use_s2() is TRUE
##
## Attaching package: 'raster'
##
##
## The following object is masked from 'package:dplyr':
##
##   select
##
## terra 1.6.17
##
## Attaching package: 'terra'
##
##
## The following object is masked from 'package:tidyr':
##
##   extract
##
## To enable caching of data, set 'options(tigris_use_cache = TRUE)'
## in your R script or .Rprofile.
##
## Attaching package: 'tigris'
##
##
## The following object is masked from 'package:terra':
```

```

##
##   blocks
##
##
## Be sure to set the download folder using 'prism_set_dl_dir()'.
##
##
## Attaching package: 'tictoc'
##
##
## The following objects are masked from 'package:terra':
##
##   shift, size
##
##
## The following object is masked from 'package:raster':
##
##   shift
##
##
## Attaching package: 'tidylog'
##
##
## The following object is masked from 'package:raster':
##
##   select
##
##
## The following objects are masked from 'package:dplyr':
##
##   add_count, add_tally, anti_join, count, distinct, distinct_all,
##   distinct_at, distinct_if, filter, filter_all, filter_at, filter_if,
##   full_join, group_by, group_by_all, group_by_at, group_by_if,
##   inner_join, left_join, mutate, mutate_all, mutate_at, mutate_if,
##   relocate, rename, rename_all, rename_at, rename_if, rename_with,
##   right_join, sample_frac, sample_n, select, select_all, select_at,
##   select_if, semi_join, slice, slice_head, slice_max, slice_min,
##   slice_sample, slice_tail, summarise, summarise_all, summarise_at,
##   summarise_if, summarize, summarize_all, summarize_at, summarize_if,
##   tally, top_frac, top_n, transmute, transmute_all, transmute_at,
##   transmute_if, ungroup
##
##
## The following objects are masked from 'package:tidyr':
##
##   drop_na, fill, gather, pivot_longer, pivot_wider, replace_na,
##   spread, uncount
##
##
## The following object is masked from 'package:stats':
##
##   filter
##

```

```

##
## Loading required package: Matrix
##
##
## Attaching package: 'Matrix'
##
##
## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack
##
##
## Loading required package: foreach
##
##
## Attaching package: 'foreach'
##
##
## The following objects are masked from 'package:purrr':
##
##   accumulate, when
##
##
## Loading required package: parallel
##
## This is INLA_22.12.16 built 2022-12-23 13:43:44 UTC.
## - See www.r-inla.org/contact-us for how to get help.
## - To enable PARDISO sparse library; see inla.pardiso()
##
##
## Attaching package: 'INLA'
##
##
## The following object is masked from 'package:fixest':
##
##   f
##
##
## Your original .Renviron will be backed up and stored in your R HOME directory if needed.
##
## Your API key has been stored in your .Renviron and can be accessed by Sys.getenv("CENSUS_API_KEY").
## To use now, restart R or run 'readRenviron("~/Renviron")'
##
## mutate: converted 'FIPS' from integer to character (0 new NA)

COD_raw <- readRDS("data_raw/age_standardized_rates/age_standardized_rates-county_month-all_sex-all_mar")

COD <-
  COD_raw %>%
  filter(year >= 2006 &
         race_eth == "all" &
         age_group == "all_ages" &
         sex == "all" &
         race_eth == "all") %>%

```

```
group_by(fipsihme, year, month) %>%
summarise(pop = sum(pop), n_deaths = sum(n_deaths), rate = n_deaths / pop) %>%
ungroup()
```

```
## filter: removed 15,051,816 rows (96%), 559,500 rows remaining
## group_by: 3 grouping variables (fipsihme, year, month)
## summarise: now 559,500 rows and 6 columns, 2 group variables remaining (fipsihme, year)
## ungroup: no grouping variables
```

```
fipsihme_sf <-
  counties(year = 2020) %>% # year should be 2006 once those data are working again
  filter(STATEFP %in% non_CONUS_FIPS == F) %>%
  dplyr::mutate(COUNTYFP = paste0(STATEFP, COUNTYFP)) %>%
  dplyr::select(COUNTYFP) %>%
  left_join(
    ihme_fips %>% dplyr::select(COUNTYFP = orig_fips, ihme_fips),
    by=c("COUNTYFP")
  ) %>%
  mutate(
    ihme_fips = ifelse(is.na(ihme_fips), COUNTYFP, ihme_fips)
  ) %>%
  dplyr::select(fipsihme = ihme_fips, COUNTYFP)
```

```
## filter: removed 126 rows (4%), 3,108 rows remaining
## left_join: added one column (ihme_fips)
##           > rows only in x   3,062
##           > rows only in y   (   31)
##           > matched rows      46
##           >                    =====
##           > rows total        3,108
## mutate: changed 3,062 values (99%) of 'ihme_fips' (3062 fewer NA)
```

```
COD_sf <-
  COD %>%
  left_join(
    fipsihme_sf %>% dplyr::select(fipsihme, geometry)
  ) %>%
  st_as_sf() #key step to recover sf geometry for some reason
```

```
## Joining, by = "fipsihme"
## left_join: added one column (geometry)
## > rows only in x 4,368
## > rows only in y ( 0)
## > matched rows 559,272 (includes duplicates)
## > =====
## > rows total 563,640
```

Maps:




Figure 1: map of death rate by county

```
acm_county_2010 <-
  COD %>%
  group_by(fipsihme, year) %>%
  summarise(annual_rate = mean(rate)) %>%
  ungroup() %>%
  filter(year == 2010) %>%
  mutate(
    annual_rate_bin = cut(annual_rate, breaks = quantile(annual_rate, seq(0, 1, .2)))
  ) %>%
  left_join(
    fipsihme_sf %>% dplyr::select(fipsihme, geometry)
  ) %>%
  st_as_sf() %>% #key step to recover sf geometry for some reason
  ggplot(aes(fill = annual_rate_bin)) +
  geom_sf(color = NA) + #color=NA to remove county borders
  scale_fill_brewer(palette = "Reds") +
  ggtitle("Death rate by county")
```

All-cause mortality rate by county for each year

```
## group_by: 2 grouping variables (fipsihme, year)
```

```
## summarise: now 46,625 rows and 3 columns, one group variable remaining (fipsihme)
```

```
## ungroup: no grouping variables
```

```
## filter: removed 43,517 rows (93%), 3,108 rows remaining
```

```
## mutate: new variable 'annual_rate_bin' (factor) with 6 unique values and <1% NA
```

```
## Joining, by = "fipsihme"
```

```
## left_join: added one column (geometry)
```

```
## > rows only in x 24
```

```
## > rows only in y ( 1)
```

```
## > matched rows 3,107 (includes duplicates)
```

```
## > =====
```


```
## > rows total 3,131
```

```
#ggsave(filename = "plots/2-17-23_COD_EDA/2010_county_deathrates.jpg", plot = acm_county_2010, device =
```

```
# library(tmap)
# tmap_test <-
#   COD %>%
#   group_by(fipsihme, year) %>%
#   summarise(annual_rate = mean(rate)) %>%
#   ungroup() %>%
#   filter(year == 2010) %>%
#   left_join(
#     fipsihme_sf %>% dplyr::select(fipsihme, geometry)
#   ) %>%
```

```
#   st_as_sf() #key step to recover sf geometry for some reason
#
# tmap_test %>%
#   dplyr::filter(!sf::st_is_empty(geometry)) %>%
#   tm_shape() +
#   tm_fill("annual_rate")
#
# tmap_save(tmap_test, filename = "plots/tmap.jpg")
```

Time series



Whole US Should match readily available data

```

### For some reason this code chunk is terminating R, I have no clue why
# monthly_US_deaths <-
#   COD_sf %>%
#   mutate(yearmonth = paste(year, month, sep="-")) %>%
#   mutate(yearmonth = ym(yearmonth)) %>%
#   group_by(yearmonth) %>%
#   summarise(n_deaths = sum(n_deaths), us_pop = sum(pop)) %>%
#   ggplot(aes(x = yearmonth, y = n_deaths)) +
#   geom_line() +
#   ggtitle("Deaths by month in the US, 2006-2020") +
#   xlab("Year") + ylab("Total deaths")

#ggsave(filename = "plots/2-17-23_COD_EDA/monthly_US_deaths.jpg", plot=monthly_US_deaths, device="jpg")

```

```

largest_counties <-
  COD %>% ungroup() %>%
  filter(year == 2010) %>% #could have chosen different year
  dplyr::select(fipsihme, pop, year) %>%
  distinct() %>%
  arrange(desc(pop)) %>%
  slice(1:10) %>%
  left_join(
    counties() %>% dplyr::select(GEOID, NAME),
    by=c("fipsihme" = "GEOID")
  ) %>%
  select(-c(pop, year)) %>%
  st_drop_geometry()

```

Time series of ACM rate and raw death counts for 10 largest counties

```

## ungroup: no grouping variables

## filter: removed 522,204 rows (93%), 37,296 rows remaining

## distinct: removed 34,188 rows (92%), 3,108 rows remaining

## slice: removed 3,098 rows (>99%), 10 rows remaining

## Retrieving data for the year 2020

## left_join: added 2 columns (NAME, geometry)

##           > rows only in x           0

##           > rows only in y   (3,224)

##           > matched rows          10

```



```

##           >           =====

##           > rows total           10

## select: dropped 2 variables (pop, year)

monthly_deaths_largest_counties <-
  COD_sf %>%
  right_join(largest_counties) %>%
  mutate(yearmonth = paste(year, month, sep="-")) %>%
  mutate(yearmonth = ym(yearmonth)) %>%
  ggplot(aes(x = yearmonth, y = rate)) +
  geom_line(aes(color = NAME))

## Joining, by = c("fipsihme", "geometry")

## right_join: added one column (NAME)

##           > rows only in x   (561,840)

##           > rows only in y           0

##           > matched rows       1,800   (includes duplicates)

##           >           =====

##           > rows total           1,800


## mutate: new variable 'yearmonth' (character) with 180 unique values and 0% NA

## mutate: converted 'yearmonth' from character to Date (0 new NA)

monthly_deaths_largest_counties +
  ggtitle("Death rate in largest counties, 2006-2020") +
  xlab("Year") + ylab("Death rate (per 100,000)")

```

Death rate in largest counties, 2006–2020



```
# counties to look closely at
large_counties =
  list("Los Angeles County" = "06037",
        "Harris County" = "48201",
        "Miami-Dade County" = "12086",
        "Cook County" = "17031",
        "New York County" = "36061",
        "San Francisco County" = "06075")

medium_counties =
  list("Jefferson County" = "01073",
        "Mecklenburg County" = "37119",
        "Lucas County" = "39095",
        "Fresno County" = "06019",
        "Boulder County" = "08013",
        "East Baton Rouge Parish" = "22033")

small_counties =
  list("Stearns County" = "27145",
        "Pike County" = "01109",
        "Clark County" = "18019",
        "Steuben County" = "36101",
        "Aroostook County" = "23003",
        "Douglas County" = "41019")

# Extract data from the lists above, make label for size category
```

```

pull_county_data_by_size = function(size) {
  COD %>%
    filter(fipsihme %in% unname(unlist(!sym(glue::glue("{size}_counties"))))) %>%
    mutate(type = size)
}

county_specific_df <- lapply(c("small", "medium", "large"), pull_county_data_by_size) %>% bind_rows()

## filter: removed 558,420 rows (>99%), 1,080 rows remaining

## mutate: new variable 'type' (character) with one unique value and 0% NA

## filter: removed 558,600 rows (>99%), 900 rows remaining

## mutate: new variable 'type' (character) with one unique value and 0% NA

## filter: removed 558,420 rows (>99%), 1,080 rows remaining

## mutate: new variable 'type' (character) with one unique value and 0% NA

monthly_rates_largecounty <-
  county_specific_df %>%
  filter(type == "large") %>%
  left_join(
    counties() %>% dplyr::select(GEOID, NAME),
    by=c("fipsihme" = "GEOID")
  ) %>%
  mutate(yearmonth = paste(year, month, sep="-")) %>%
  mutate(yearmonth = ym(yearmonth)) %>%
  ggplot(aes(x = yearmonth, y = rate)) +
  geom_line(aes(color = NAME))

## filter: removed 1,980 rows (65%), 1,080 rows remaining

## Retrieving data for the year 2020

## left_join: added 2 columns (NAME, geometry)

##           > rows only in x           0

##           > rows only in y   (3,228)

##           > matched rows       1,080

##           >           =====

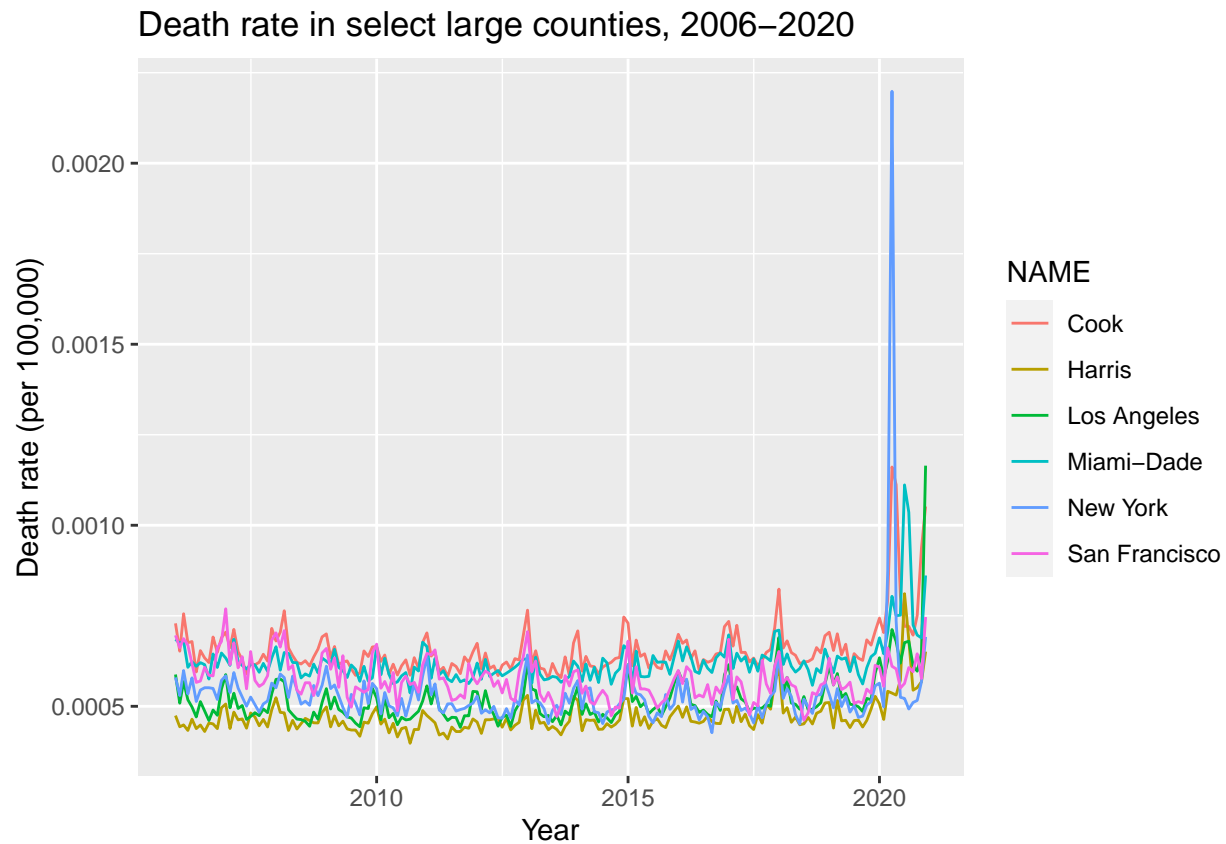
##           > rows total         1,080

## mutate: new variable 'yearmonth' (character) with 180 unique values and 0% NA

## mutate: converted 'yearmonth' from character to Date (0 new NA)

```

```
monthly_rates_largecounty +
  ggtitle("Death rate in select large counties, 2006-2020") +
  xlab("Year") + ylab("Death rate (per 100,000)")
```



```
monthly_rates_mediumcounty <-
  county_specific_df %>%
  filter(type == "medium") %>%
  left_join(
    counties() %>% dplyr::select(GEOID, NAME),
    by=c("fipsihme" = "GEOID")
  ) %>%
  mutate(yearmonth = paste(year, month, sep="-")) %>%
  mutate(yearmonth = ym(yearmonth)) %>%
  ggplot(aes(x = yearmonth, y = rate)) +
  geom_line(aes(color = NAME))
```

```
## filter: removed 2,160 rows (71%), 900 rows remaining
```

```
## Retrieving data for the year 2020
```

```
## left_join: added 2 columns (NAME, geometry)
```

```
## > rows only in x 0
```

```
##           > rows only in y  (3,229)

##           > matched rows      900

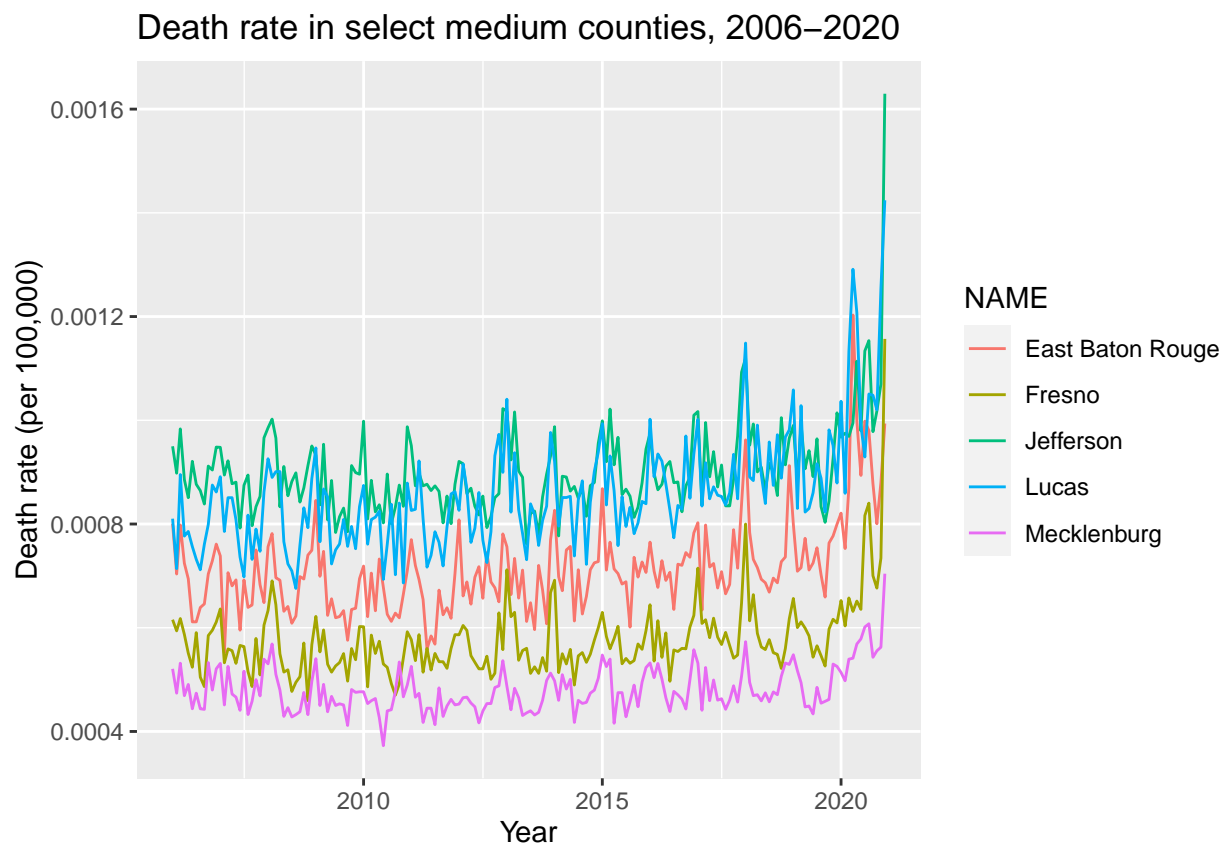
##           >                =====

##           > rows total        900

## mutate: new variable 'yearmonth' (character) with 180 unique values and 0% NA

## mutate: converted 'yearmonth' from character to Date (0 new NA)

monthly_rates_mediumcounty +
  ggtitle("Death rate in select medium counties, 2006-2020") +
  xlab("Year") + ylab("Death rate (per 100,000)")
```



```
monthly_rates_smallcounty <-
  county_specific_df %>%
  filter(type == "small") %>%
  left_join(
    counties() %>% dplyr::select(GEOID, NAME),
    by=c("fipsihme" = "GEOID")
  ) %>%
  mutate(yearmonth = paste(year, month, sep="-")) %>%
```

```
mutate(yearmonth = ym(yearmonth)) %>%
  ggplot(aes(x = yearmonth, y = rate)) +
  geom_line(aes(color = NAME))
```

```
## filter: removed 1,980 rows (65%), 1,080 rows remaining
```

```
## Retrieving data for the year 2020
```

```
## left_join: added 2 columns (NAME, geometry)
```

```
##           > rows only in x           0
```

```
##           > rows only in y   (3,228)
```

```
##           > matched rows       1,080
```


```
##           >                      =====
```

```
##           > rows total         1,080
```

```
## mutate: new variable 'yearmonth' (character) with 180 unique values and 0% NA
```

```
## mutate: converted 'yearmonth' from character to Date (0 new NA)
```

```
monthly_rates_smallcounty +
  ggtitle("Death rate in select medium counties, 2006-2020") +
  xlab("Year") + ylab("Death rate (per 100,000)")
```



Sanity checks

Compare population counts to ACS5 counts Vast majority of counties appear colinear but there are some exceptions

```
acs2009to2020 <- read_csv("data/acs2009to2020.csv")
```


```
## Rows: 38646 Columns: 58
## -- Column specification -----
## Delimiter: ","
## chr (2): GEOID, NAME
## dbl (56): year, pop, male, male_u15, male_15to19, male_20to24, male_25to29, ...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
COD_ACS <-
  COD %>%
  filter(year >= 2009) %>%
  left_join(
    acs2009to2020,
    by=c("fipsihme" = "GEOID", "year")
  )
```

```
## filter: removed 111,924 rows (20%), 447,576 rows remaining
## left_join: added 57 columns (pop.x, NAME, pop.y, male, male_u15, ...)
##           > rows only in x      144
##           > rows only in y ( 1,360)
##           > matched rows      447,432
##           >                    =====
##           > rows total        447,576
```

```
COD_ACS %>%
  ggplot(aes(x = pop.x, y = pop.y)) +
  geom_point(size=.1) +
  geom_abline(slope = 1, linewidth = .1)
```


```
## Warning: Removed 144 rows containing missing values ('geom_point()').
```



Investigate non-collinear counties Set a 10% margin of error

```
COD_ACS %>%
  filter(pop.y < .9 * pop.x | pop.y > 1.1 * pop.x) %>%
  ggplot(aes(x = pop.x, y = pop.y)) +
  geom_point(size=.1) +
  geom_abline(slope = 1, linewidth = .1)
```

filter: removed 439,188 rows (98%), 8,388 rows remaining



```
pop_mismatch <-
  COD_ACS %>%
  filter(pop.y < .9 * pop.x | pop.y > 1.1 * pop.x)
```

```
## filter: removed 439,188 rows (98%), 8,388 rows remaining
```

#shows that 2009 (and that ACS grouping) has more of these mismatches

```
pop_mismatch %>%
  dplyr::select(fipsihme, year, NAME, pop.x, pop.y) %>%
  distinct() %>%
  group_by(year) %>%
  summarise(n = n())
```

```
## distinct: removed 7,689 rows (92%), 699 rows remaining
```

```
## group_by: one grouping variable (year)
```

```
## summarise: now 12 rows and 2 columns, ungrouped
```

```
## # A tibble: 12 x 2
##   year      n
##   <dbl> <int>
## 1  2009    139
## 2  2010     53
```

```
## 3 2011 50
## 4 2012 46
## 5 2013 50
## 6 2014 51
## 7 2015 49
## 8 2016 57
## 9 2017 57
## 10 2018 48
## 11 2019 50
## 12 2020 49
```

#AK and CO have a disproportionate amount of these mismatch counties

```
pop_mismatch %>%
  dplyr::select(fipsihme, year, NAME, pop.x, pop.y) %>%
  mutate(statefip = str_sub(fipsihme, 1, 2)) %>%
  left_join(STATEFP, by = c("statefip" = "FIPS")) %>%
  distinct() %>%
  group_by(state) %>%
  summarise(n = n())
```

mutate: new variable 'statefip' (character) with 34 unique values and 0% NA

left_join: added one column (state)

```
##           > rows only in x      0
##           > rows only in y (   21)
##           > matched rows      8,388
##           >                      =====
##           > rows total        8,388
```

distinct: removed 7,689 rows (92%), 699 rows remaining

group_by: one grouping variable (state)

summarise: now 34 rows and 2 columns, ungrouped

```
## # A tibble: 34 x 2
##   state      n
##   <chr> <int>
## 1 AK      63
## 2 AL       1
## 3 AZ      15
## 4 CA       2
## 5 CO      53
## 6 FL       4
## 7 GA      25
## 8 HI      12
## 9 IA       1
## 10 ID     18
## # ... with 24 more rows
## # i Use 'print(n = ...)' to see more rows
```

Counties where deaths > 50% of population Saw some crop up when doing an analysis earlier but with this cleaner process, those seem to have gone away (good)

```
COD %>%  
  filter(n_deaths >= pop * .5)
```

```
## filter: removed all rows (100%)
```

```
## # A tibble: 0 x 6  
## # ... with 6 variables: fipsihme <chr>, year <dbl>, month <dbl>, pop <dbl>,  
## #   n_deaths <dbl>, rate <dbl>  
## # i Use 'colnames()' to see all variable names
```